


- most common indexing data structure used in relational databases systems
- optimized for disk-based indexing
- minimizing disk accesses for indexing

B+ Tree is an m-way with order M

- M -> maximum number of keys in each node
- M+1 -> max children of each node

Node structure for m=3



B+ Tree Properties:



- All node (except root) must be $\frac{1}{2}$ full min
- Root node doesn't have to be $\frac{1}{2}$ full
- Insertions are always done at leaves
- Leaves are stored as a DLL
- Keys in nodes are kept sorted

Insert 42, 29, 81, Into M=3



Insert 99, 35, 2



Now two nodes and no way to
get between them



root



Copy smallest value of new node
(which is 81)




- Internal nodes
 - Store pointers to children and keys
 - Acts as the index
 - Just indexes relevant info -> pack keys so minimize depth
- Leaf nodes
 - Store keys and also data

B+ Tree is extension of B tree, but in B tree internal nodes store keys and values

- B+ tree minimizes disk access (disk space indexing)
- B tree for in memory indexing


Insert 2




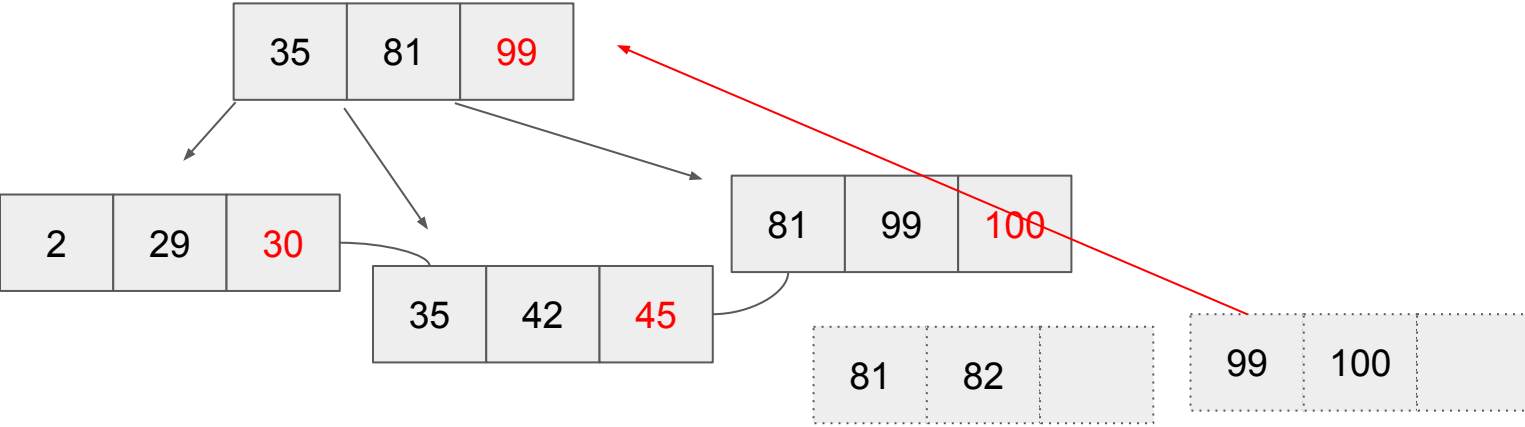
Split 2, 29, 35, 42 in half into diff nodes




Operations to split nodes and insert
is still faster than doing more disk
accesses



Insert: 100, 30, 45, 82



Insert 4



When split internal node, gets moved up to parent not copied